



SIMONE BONFRATE

Enhancing Predictive Attack Analysis with GANs

New Frontiers in Cybersecurity

Table of contents

01 The Evolving Threat Landscape

03 GANs in Cybersecurity

05 Intrusion detection

07 Limitations and Ethical Considerations

02 Generative Adversarial Networks

04 Malware detection

06 Anomaly detection

08 Conclusions

The Evolving Threat Landscape

Advanced Persistent Threats (APTs):

Highly targeted and persistent attacks designed to infiltrate systems undetected for extended periods.

Ransomware

Malware that encrypts data and demands a ransom for its release. Ransomware attacks are becoming increasingly prevalent and sophisticated.

Polymorphic Malware

Malware that constantly changes its code to evade detection.

AI-powered Attacks

Attacks leveraging machine learning and AI to bypass security systems and automate attacks.


Data Breaches

Large-scale data breaches continue to compromise sensitive personal and corporate information.

Generative Adversarial Networks

Generative Adversarial Networks (GANs) are a powerful class of machine learning models capable of generating new data instances that resemble the training data.

- **Generator:** This network tries to create synthetic data that mimics the real data. It learns to map random noise into realistic data samples.
- **Discriminator:** This network tries to distinguish between real data and the synthetic data generated by the generator. It acts as a critic, evaluating the quality of the generated data.



Generative Adversarial Networks (GANs): Key Features



Adversarial Training

The competitive training process drives the generator to produce increasingly realistic outputs.



Unconstrained Generation

GANs can generate complex data instances beyond simple interpolation of the training data.



High-Quality Data

GANs are capable of generating realistic and high-fidelity synthetic data.

Generative Adversarial Networks: Architectures evolutions

While the basic GAN architecture is powerful, several variations have been developed to address specific limitations and improve performance. The choice of GAN architecture depends on the specific application and dataset characteristics. Each type offers unique advantages and disadvantages regarding training stability, control over the generated data, and suitability for different data types.

Wasserstein GAN

Addresses training instability issues of the Vanilla GAN by using the Wasserstein distance as a loss function. Generally more stable than Vanilla GANs

Conditional GAN

Allows for controlled generation of data by conditioning the generator on specific labels or inputs. This gives more control over the generated data.

Table GAN

Specifically designed for generating tabular data, which is particularly relevant for cybersecurity datasets that often involve structured data like network logs or system events.

GANs in Cybersecurity: Key Applications

Thanks to their capacity for generating high-quality data, GANs are particularly helpful in addressing dataset imbalances, as they can easily generate synthetic examples of infrequent attacks.

Malware Detection

GANs can generate synthetic malware samples to augment training datasets for machine learning models, improving their ability to detect both known and unknown malware variants

Intrusion Detection

For training Intrusion Detection Systems (IDS) which make accurate detection of malicious activities.

Anomaly Detection

For modelling normal system behavior and identify deviations from this norm, flagging potential security incidents.


Malware Detection: LSTM Meets GAN

This research explores the use of GANs in conjunction with Long Short-Term Memory (LSTM) networks for malware detection. LSTMs are particularly well-suited for analyzing sequential data, such as the byte sequences found in malware executables.

Approach:

- 1. Data Augmentation:** A GAN is trained to generate synthetic malware samples. This aims to increase the size and diversity of the training dataset, addressing the limitations of real-world malware datasets which are often scarce and may not capture the full range of malicious behavior.
- 2. Feature Extraction:** The generated and real malware samples are processed to extract relevant features, such as opcodes, function calls, and API calls.
- 3. Malware Classification:** An LSTM network is trained to classify these features as either benign or malicious. The combination of real and synthetic data helps the LSTM to learn more robust patterns.

This approach leverages the strengths of both GANs and LSTMs to create a more effective malware detection system.




Ishita Gupta et al. 2024, Leveraging LSTM and GAN for Modern Malware Detection

Malware Detection: LSTM Meets GAN

The model was evaluated using the VirusShare dataset, a large and diverse collection of malware samples and it has been obtained the follow key results:

1. Accuracy: The LSTM-GAN model achieved an accuracy of 98.82%.
2. Comparison to traditional methods: This significantly outperforms traditional classifiers such as Random Forest, SVM, and Decision Trees.
3. Generalization: The inclusion of GAN-generated data improved the model's ability to generalize to unseen malware samples.

This high accuracy demonstrates the potential of GANs to enhance the effectiveness of malware detection systems.




Intrusion Detection: GAN-Enhanced Network Intrusion Detection Systems (NIDS)

NIDS often struggle with imbalanced datasets, where normal network traffic significantly outweighs malicious activities. GANs can address this issue by generating synthetic samples of infrequent attacks, leading to more robust model training.

In these research, Vanilla GAN, WGAN and CTGAN had been tested for detecting which architecture is the best for augmenting CIC-IDS2017 dataset. . Careful evaluation was performed to ensure the generated data's quality and similarity to real-world patterns.

Next, a random forest classifier was trained using the augmented dataset. This allows the model to better learn the characteristics of both normal and anomalous network activities, improving detection accuracy.



Xinxing Zhao et al. 2024, Enhancing network intrusion detection performance using generative adversarial networks

Intrusion Detection: GAN-Enhanced Network Intrusion Detection Systems (NIDS)

The integration of GAN-generated data significantly improved the performance of the Random Forest classifier in detecting Botnet attacks within the CIC-IDS2017 dataset.

Key Findings:


1. Vanilla GAN and WGAN demonstrated superior performance compared to CTGAN in similarity respect to dataset and achieving perfect precision, recall, and F1-score (1.00) in detecting Botnet attacks even with a relatively small augmentation of the dataset (4x).
2. Dataset Augmentation improves models performance using all GAN models, thanks to GAN ability to generalize to unseen attacks, making the NIDS more robust.

Features	GAN Models		
	Vanilla GAN	WGAN	CTGAN
Feature 1	0.8979	0.8986	0.8295
Feature 2	0.7608	0.7607	0.5531
Feature 3	0.9673	0.9653	0.9147
Feature 4	0.9013	0.8977	0.8316
Feature 5	0.8962	0.8956	0.8520
Feature 6	0.9669	0.9655	0.8958
Feature 7	0.8750	0.8750	0.7001
Feature 8	0.9999	0.9999	0.9997

Model (Numbers)	Generated Samples			Original Samples		
	Precision	Recall	F1-Score	Precision	Recall	F1-Score
WGAN (4 times)	0.97	1.00	0.98	1.0	0.74	0.85
WGAN (49 times)	1.00	1.00	1.00	1.0	0.76	0.87
WGAN (99 times)	1.00	1.00	1.00	1.0	0.82	0.90
Vanilla GAN (4 times)	0.98	0.99	0.98	1.0	0.66	0.80
Vanilla GAN (49 times)	1.00	1.00	1.00	1.0	0.76	0.87
Vanilla GAN (99 times)	1.00	1.00	1.00	1.0	0.81	0.90
CTGAN (4 times)	0.99	0.68	0.81	1.0	0.43	0.60
CTGAN (49 times)	0.97	0.92	0.95	1.0	0.76	0.86
CTGAN (99 times)	0.95	0.99	0.97	1.0	0.77	0.87

Anomaly Detection: The Attention-GAN Approach

This research introduces Attention-GAN, a novel architecture combining the generative capabilities of GANs with the focused feature extraction of Convolutional Neural Networks (CNNs) enhanced by attention mechanisms. This approach aims to improve anomaly detection in cybersecurity by focusing on the most relevant features within network traffic data.



Mohammed Abo Sen et al. 2024, Attention-GAN for Anomaly Detection: A Cutting-Edge Approach to Cybersecurity Threat Management

Anomaly Detection: The Attention-GAN Approach

The results obtained on CICIDS2017 and KDD demonstrate an accuracy higher than 97% (its effectiveness in intrusion detection (97.93% on CICIDS2017 dataset and 99.69% on KDD dataset)).

Direct comparisons with other state-of-the-art methods show that Attention-GAN significantly outperforms the alternatives, demonstrating an effective integration of attention mechanisms and GAN for cybersecurity.

Model / Method	F1-Score
Ensemble Methods & Traditional Machine Learning Chakir et al., 2023	0.90
Support Vector Machine (SVM) Vijayanand et al., 2017	0.89
Decision Tree Wei et al., 2010	0.88
Anomaly Detection (Hardware Performance Counters) Krishnamurthy et al., 2020	0.86
Signature-Based & Behavior-Based Detection Visoottiviseth et al., 2020	0.85
Deep Learning Techniques (Baseline) Salam et al., 2023	0.94
Attention-GAN (Proposed Method)	>0.97

GANs technical limitations

Computational Cost

could be misused to generate sophisticated malware, synthetic phishing emails, or deepfakes for social engineering attacks, exacerbating existing security threats.

Data Dependency

The performance of GANs heavily relies on the quality and quantity of the training data. Insufficient or biased data can lead to poor model performance.

Interpretability

If the training data contains biases, the GAN may amplify those biases in the generated data, leading to unfair or discriminatory outcomes.

Model Collapse

A common problem where the generator produces only a limited variety of outputs, hindering its ability to generate diverse and realistic samples.

Training Instability

GAN training can be notoriously unstable, requiring careful hyperparameter tuning and the selection of appropriate architectures.

GANs ethical considerations

Malicious Use

Training GANs can be computationally expensive and time-consuming, requiring substantial computing resources.

Privacy Violation

GANs could generate potential privacy violations if the generated data is too similar to real data and could be used to re-identify individuals.

Bias Amplification

Understanding the decision-making process of GANs can be challenging, making it difficult to assess the reliability of the generated data and the model's predictions.

Lack of Transparency and Explainability

The nature of GANs can make it difficult to understand their decision-making processes, hindering accountability and trust.

Misinformation and Disinformation

Responsible development and deployment of GANs require careful consideration of these ethical implications

Conclusions and Future Works

This research demonstrates the significant potential of GANs to enhance cybersecurity defenses.

- 1. Improved Detection Accuracy:** GANs significantly improve the accuracy of malware detection and intrusion detection systems by augmenting training datasets with realistic synthetic data. Models achieved accuracy exceeding 98% in some cases.
- 2. Enhanced Generalization:** GAN-augmented models show improved ability to generalize to unseen attacks and malware variants, leading to more robust and resilient security systems.
- 3. Addressing Data Imbalance:** GANs effectively address the problem of imbalanced datasets, common in cybersecurity, by generating synthetic samples of underrepresented attack types.
- 4. Simulation of Realistic Attacks:** GANs provide a powerful tool for simulating realistic attack scenarios, enabling more effective testing and evaluation of security measures.



Conclusions and Future Works

1. **Advanced GAN Architectures:** Exploring more advanced GAN architectures, such as improved Wasserstein GANs (WGAN-GP) and other variations, to further enhance training stability and the quality of generated data.
2. **Multimodal Attacks:** Developing GANs capable of generating multimodal attacks, integrating various data types (e.g., network traffic, system logs, code) to simulate more realistic and complex threat scenarios.
3. **Explainable AI (XAI):** Developing methods to enhance the interpretability of GAN models, improving understanding of their decision-making processes and increasing trust in their predictions.
4. **Federated Learning:** Exploring federated learning approaches to train GANs on decentralized datasets from multiple organizations, preserving data privacy while still benefiting from the combined data.
5. **Adaptive Security Systems:** Integrating GANs into adaptive security systems that can automatically adjust their defenses based on the evolving threat landscape.





Thanks for your attention !

CREDITS: This presentation template was created by [Slidesgo](#), and includes icons by [Flaticon](#), and infographics & images by [Freepik](#)